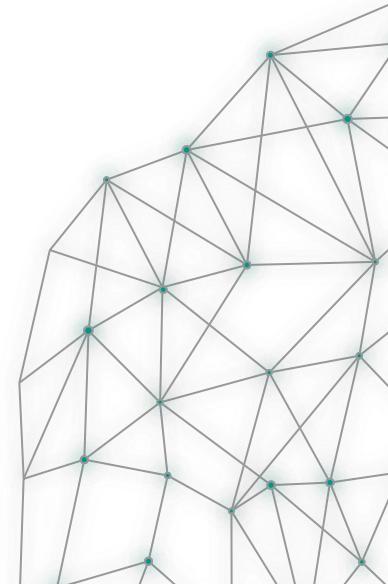


LLMs RAG & Graph Knowledge

Mohaddeseh Tabrizian

Engineering with GenAI
09.12.2025



Agenda

- Introduction and Overview RAG
- Information Retrieval and Search Foundations
- Information Retrieval with Vector Databases
- LLMs and Text Generation
- Knowledge Graphs

Introduction and Overview RAG

Understanding RAG

LLMs are powerful

- ✓ Summarize text
- ✓ Generate code
- ✓ Rewrite content

but RAG further improves them



+



LLM

New
information

Two Steps for Answering Questions

**Why Kaisers-
lautern is so
quiet on
weekends?**



Retrieval

**Why Kaisers-
lautern is so
crowded
today?**

Collect
information



Generation

Synthesize
research

Just put it in the prompt

User prompt

Why Kaiserslautem
is so crowded today?

+

News reports

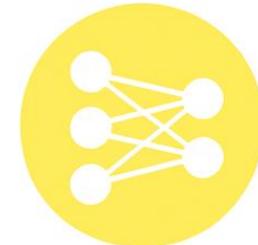
Forum posts



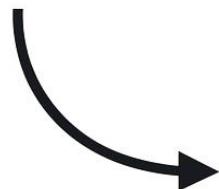
LLM

User prompt

Why Kaiserslautern
is so crowded today?



LLM



Augmentation



FCK has a game with
Dortmund today.

Retriever

- Manages knowledge base of trusted information
- Finds the most relevant information and shares information with the LLM
- Improves generation

Retriever



Knowledge base

Retreival Augmented Generation

RAG System

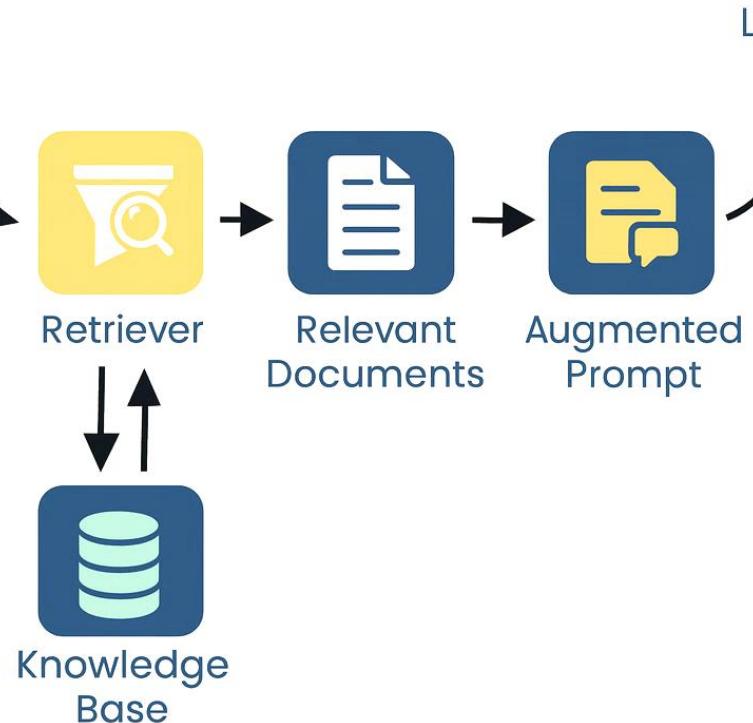
Normal
LLM Use



NORMAL LLM USE



RAG SYSTEM



Why LLMs Hallucinate



LLMs generate probable word sequences

LLMs produce word sequences based on training patterns



Knowledge gaps cause inaccurate responses

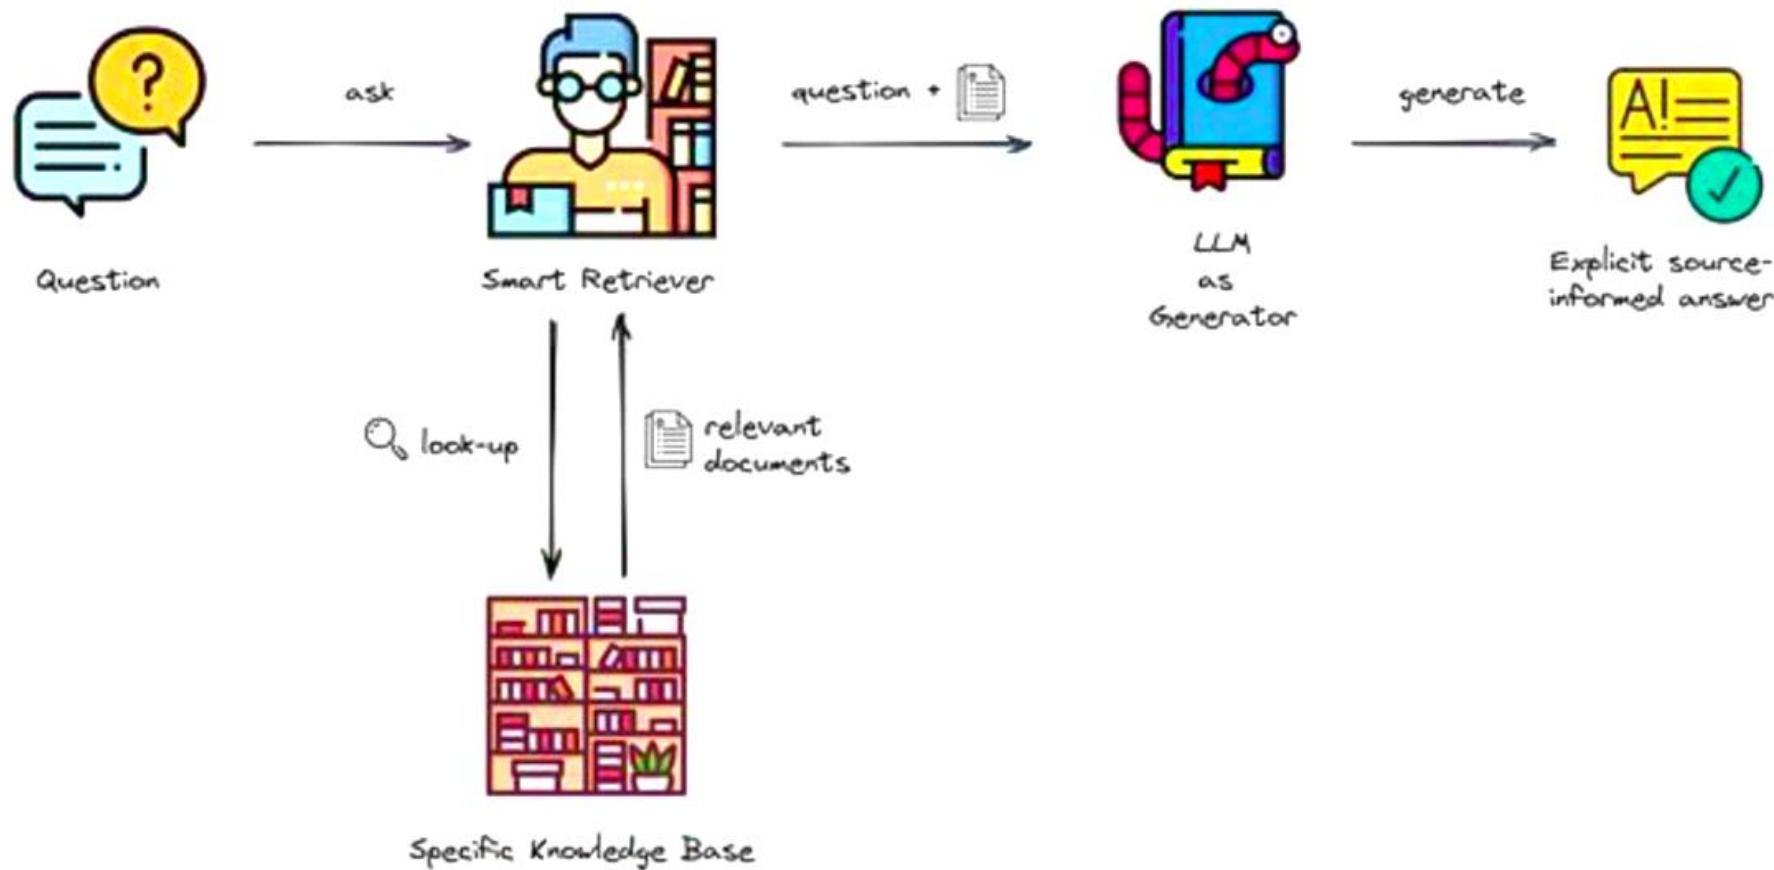
Responses might seem correct but are not



Truthful ≠ probable

LLMs are not designed to generate “truthful” text

How RAG solves the problem?



Information Retrieval and Search Foundations

Two Search Approaches



Keyword Search

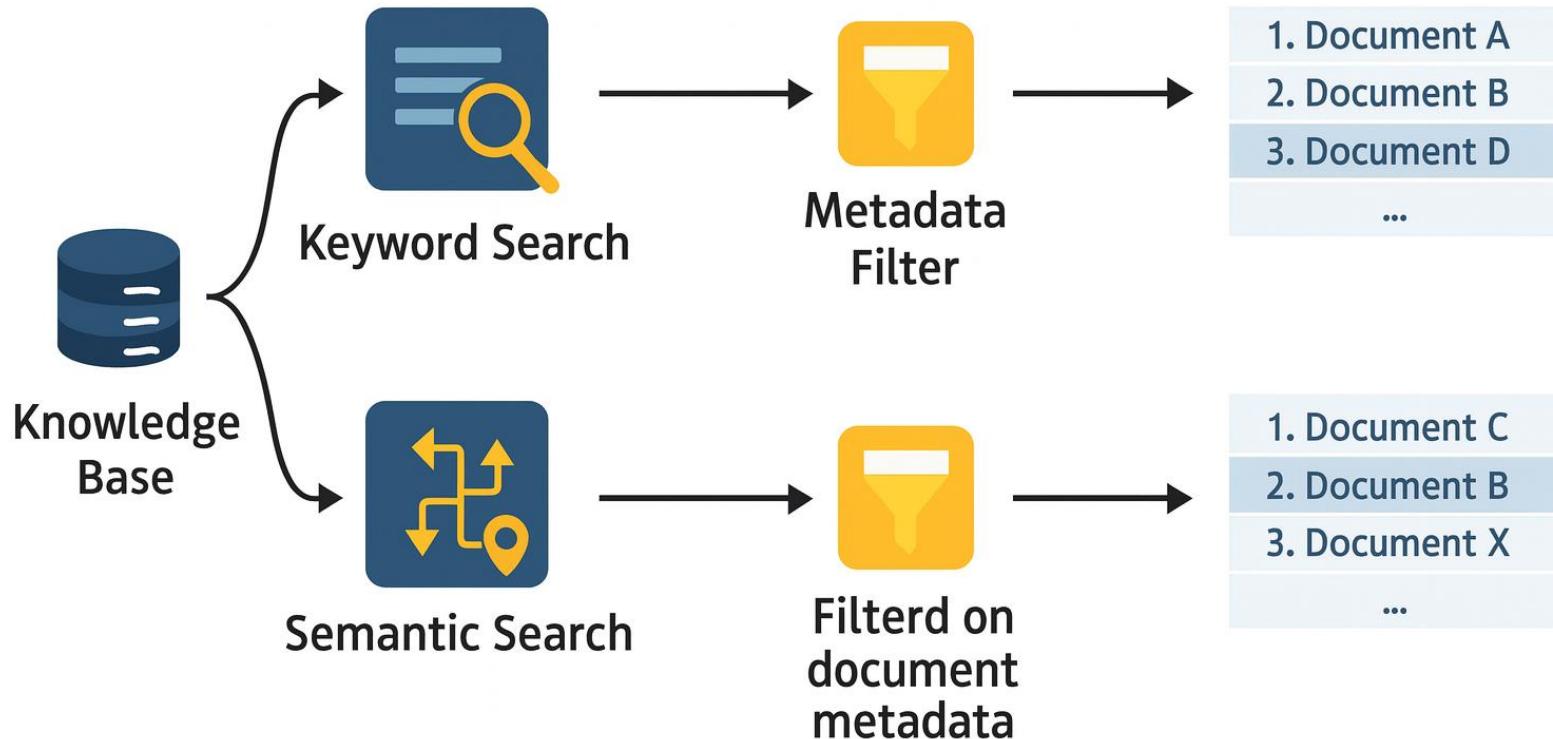
Looks for documents containing the **exact words** found in the prompt.



Semantic Search

Looks for documents with **similar meaning** to the prompt.

Search Techniques



Metadata Filtering



The Daily Maple



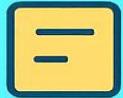
Title



Publication Date



Author



Section



Tags

```
SELECT * FROM articles
```

```
WHERE publication_date = '2023-10-01';
```

Return all articles published on specific

```
section = "Opinion"
```

```
author = "Michael Chen"
```

```
date = June to July 2024
```



How can I make
New York style
pizza at home?



**Keyword
Search**



Doc 1



Doc 2



Doc 3

Use bread flour for New
York pizza dough.

New York pizzerias stay
open late for home
delivery.

At New York Pizza Mexico,
they serve pizza

Bag of Words

Word order is ignored, only word presence and frequency matter

Prompt

Making pizza without a pizza oven

Keywords

pizza

2

without

1

making

1

oven

1

a

1

Bag of words

Sparse Vectors

Making pizza without a pizza oven

making
1

pizza
1

without
1

drink
0

pan
0

burger
0

taco
0

with
0

salad
0

tea
0

Most words aren't used. The bag of words is sparse, with few non-zero entries.

MAKING PIZZA WITHOUT A PIZZA OVEN: DOCUMENT ANALYSIS

| | 5/5 Doc 1 | 3/5 Doc 2 | 1/5 Doc 3 | 4/5 Doc 4 | 4/5 Doc 5 |
|---------|--------------|--------------|--------------|--------------|--------------|
| making | +1 | | | +1 | +1 |
| pizza | +1 | | +1 | | |
| without | +1 | +1 | | +1 | +1 |
| a | +1 | +1 | | +1 | +1 |
| oven | +1 | +1 | | +1 | +1 |

Frequency Based Scoring

Document 1

Homemade pizza in oven is better than frozen pizza

Contains: Pizza (2x) Oven (1x)

Simple Scoring = 2 points

TF Scoring = 4 points

Document 2

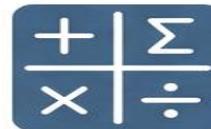
Wood-fired oven is a better oven than a stone oven for pizza

Contains: Pizza (1x) Oven (3x)

Simple Scoring = 2 points

TF Scoring = 4 points

TF-IDF

| PROBLEM | SOLUTION | FORMULA |
|--|--|--|
| <p>Basic TF scoring treats all words equally, whether they're common filler words or rare, meaningful terms.</p>  | <p>Weight terms using 'inverse document frequency' (IDF).</p>  | <p>Score = $\text{TF}(\text{word}, \text{doc}) \times \log(\text{Total docs} / \text{Docs containing word})$</p>  |

TF-IDF: Document Frequency Analysis

COUNT DOCUMENTS (DF)



Pizza: Appears in 5 out of 100 documents.

$$DF = 5/100 \quad 0.05$$

FLIP TO REWARD RARE WORDS (IDF)



High Score

$$IDF = 1 / 0.05$$

20

APPLY LOG (LOG IDF)



Still higher

$$\log(20)$$

1.30

T

The: Appears in all 100 documents.

$$DF = 100/100$$

1.0

T

Low Score

$$IDF = 1 / 1.0$$

1.0

T

Too common, no weight

$$\log(1)$$

0

Evolution of Information Retrieval Models

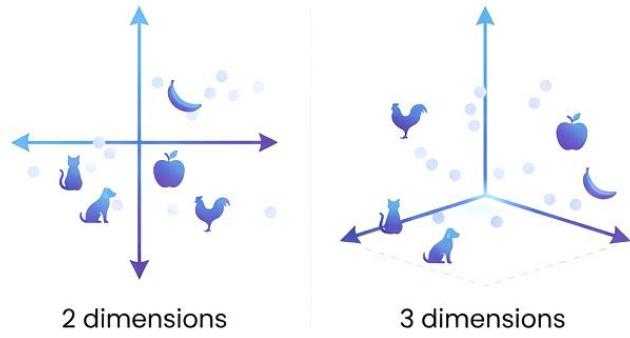


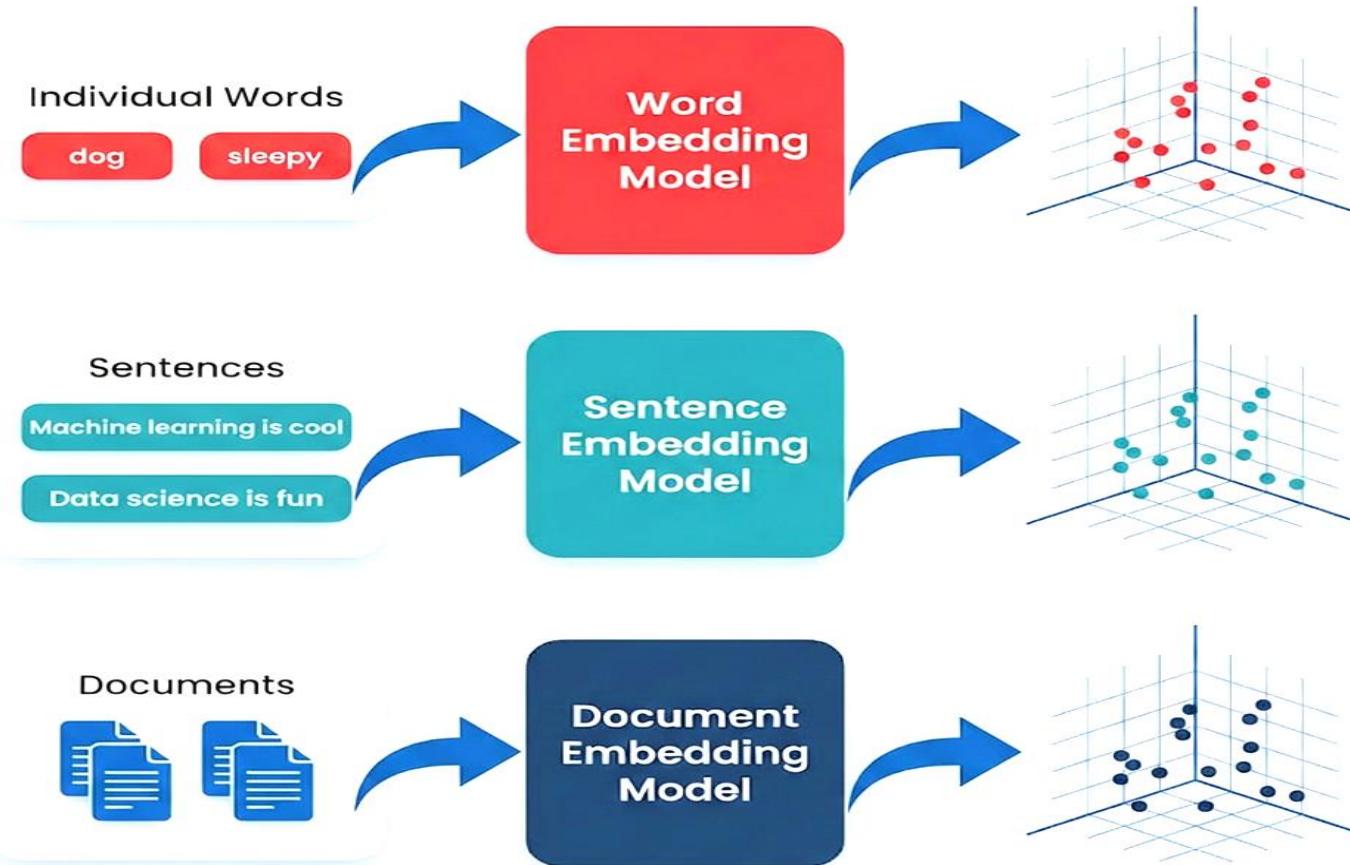
Modern systems use a slightly refined version called BM25

Semantic Search vs. Keyword Search

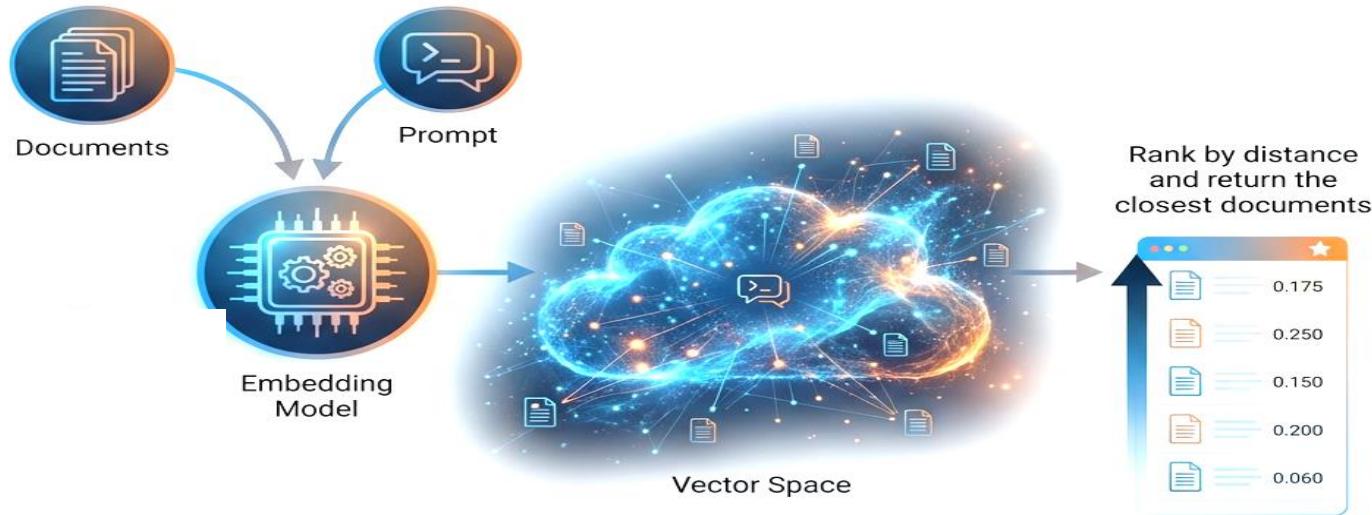
- Prompt and documents each get a vector
- Vectors compared to generate scores
- The main difference is how vectors are assigned
 - **Keyword Search:** count words
 - **Semantic Search:** use embedding model

Understanding Embedding Models



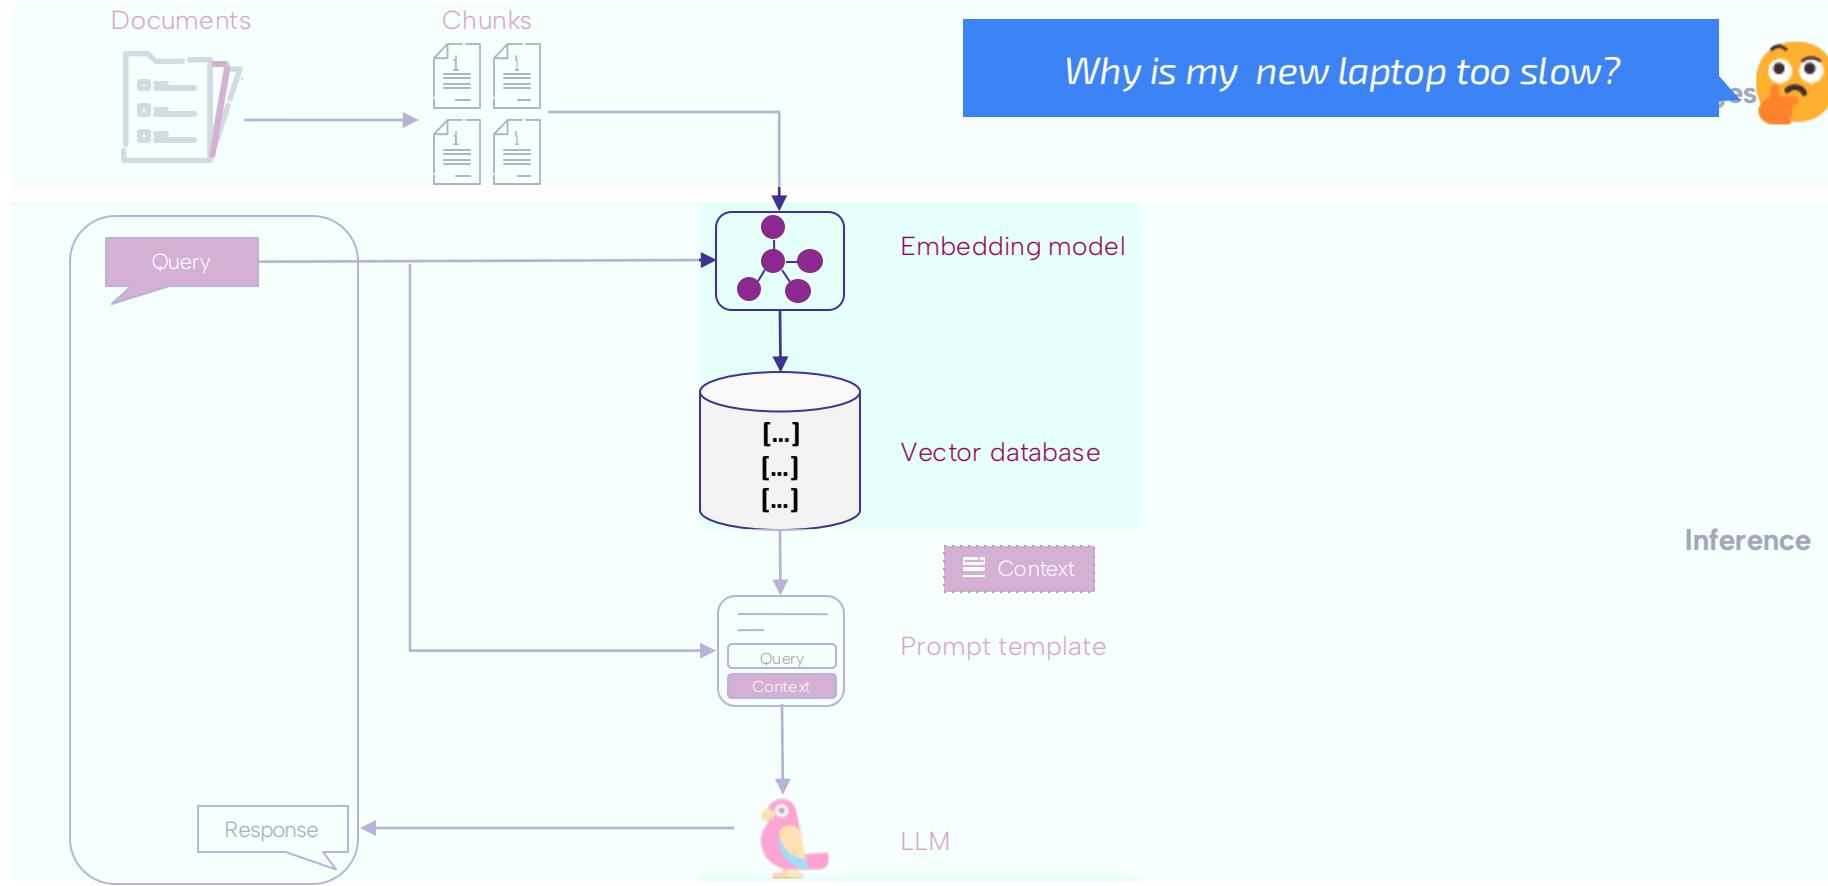


Semantic Search



Information Retrieval with Vector Databases

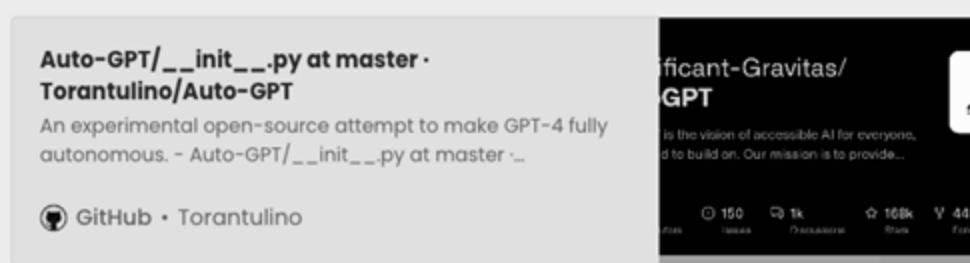
How does RAG work? | VectorDBs



Vector Databases

The Vector DB: An Overkill Solution

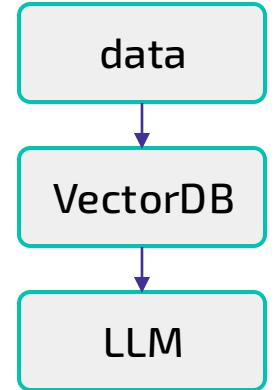
Auto-GPT relies on vector databases for faster k-nearest neighbor (kNN) searches. These databases retrieve earlier chains of thoughts, incorporating them into the current query context for GPT to provide a memory effect. However, given Auto-GPT's constraints and limitations, this approach has been criticized as excessive and unnecessarily resource-intensive.



The main argument against using vector databases stems from the cost constraints associated with Auto-GPT's chain of thoughts. A 50-step chain of thoughts would cost 14.4 dollars, while a 1000-step chain would cost significantly more. As a result, the memory size or length of the chain of thought rarely exceeds four digits. In such cases, an exhaustive search (i.e., the dot product between a 256-dim vector and a 10,000 x 256 matrix) for the nearest neighbor proves to be sufficiently efficient, taking less than one second to complete. In contrast, each GPT-4 call takes approximately 10 seconds to process, making the system GPT-bound rather than database-bound.

Vector Databases

- What is a vector DB?
- How does it work?
- Use cases



Vector Databases | Why

unstructured data

> 80%



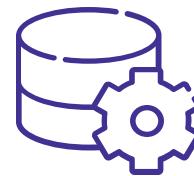
images



video



audio



LLMs and Text generation

Controlling LLM Output:

Key Techniques



Greedy Decoding (Deterministic)

Always picks highest-probability token.
Good for code/debugging, but can be bland or repetitive.



Temperature (Adjusts Sharpness)

Lower → Predictable, focused.
Higher → Varied, creative.
(Temp 0 = Greedy).



Sampling Methods (Refines Choices)

Top-k: Restricts to "k" most likely tokens.
Top-p (Nucleus): Includes tokens with cumulative probability exceeding a threshold (e.g., 90%). Adapts to certainty.



Takeaway: Tuning temperature, top-k, top-p, and penalties shapes predictability vs. creativity. Experiment to find the right fit for your application.

Choosing the Right LLM for Your RAG Application: Key Factors

Quantifiable Metrics (Performance & Cost)



Model Size (Parameters): Billions (e.g., 1-10B vs. 100-500B+). Trade-off: Capability vs. Cost.



Cost: Price per million tokens (Input/Output). Newer/larger models cost more.



Context Window: Max tokens (Prompt + Completion). More flexibility, pay per token.



Speed & Latency: Time to first token, Tokens/sec. Critical for real-time.



Training Cutoff: Knowledge cutoff date. Later is preferable for recent events.

Quality & Benchmarking (Evaluation)

Quality is hard to quantify (reasoning, readability). **Benchmarks** help compare.



Takeaway: Balance these quantifiable metrics and benchmark results against your project's specific speed, quality, and budget requirements.

Advanced Prompt Engineering for RAG: Techniques & Application

In-Context Learning (Grounding Responses)



Concept: Provide examples to guide LLM structure, tone, and output format.

Types: One-shot (single example), Few-shot (multiple examples).

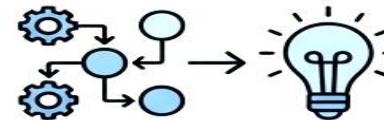
Implementation: Hard-code stable examples or use RAG to retrieve contextually relevant past examples.

Benefit: Grounds the LLM, improving response quality and consistency.



RAG for Examples: Retrieve successful past interactions to inject dynamic, high-quality examples.

Reasoning & Chain of Thought (Step-by-Step Logic)

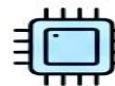


Concept: Encourage the LLM to plan and reason before generating a final answer.

Techniques: 'Think step-by-step', 'Scratchpad' tags (intermediate thinking space), Chain of Thought (generate steps, then follow them).

Benefit: Increases accuracy for complex tasks (math, coding, planning), makes errors traceable.

Evolution: Rise of 'Reasoning Models' with built-in planning & thought-process tokens.



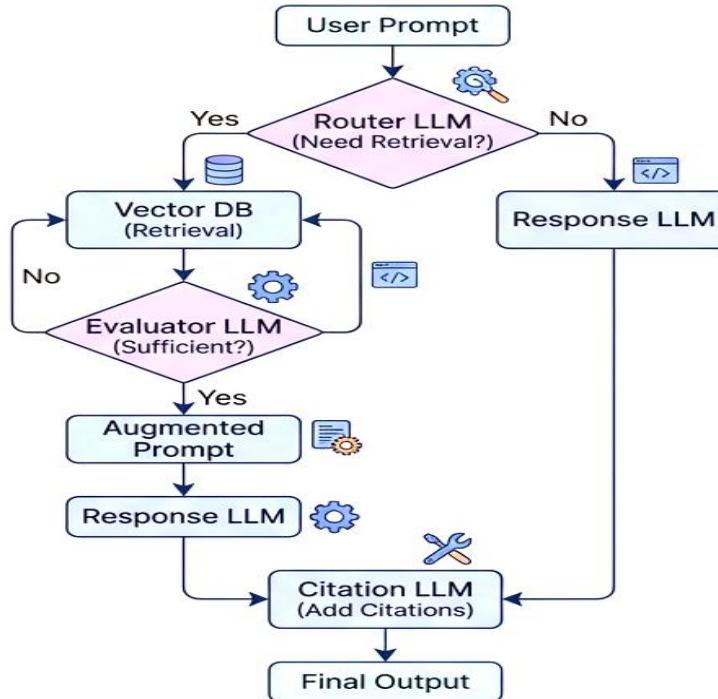
Reasoning Models: Designed specifically for multi-step complex workflows right out of the box.



Goal: Leverage these techniques to produce more accurate, structured, and high-quality RAG outputs.

Agentic Workflows in RAG: Modular & Specialized Systems

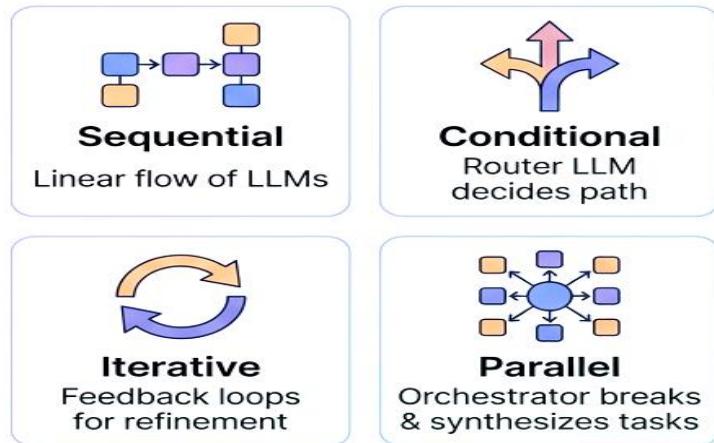
Example Agentic RAG System (Multi-Step Process)



Key Concepts

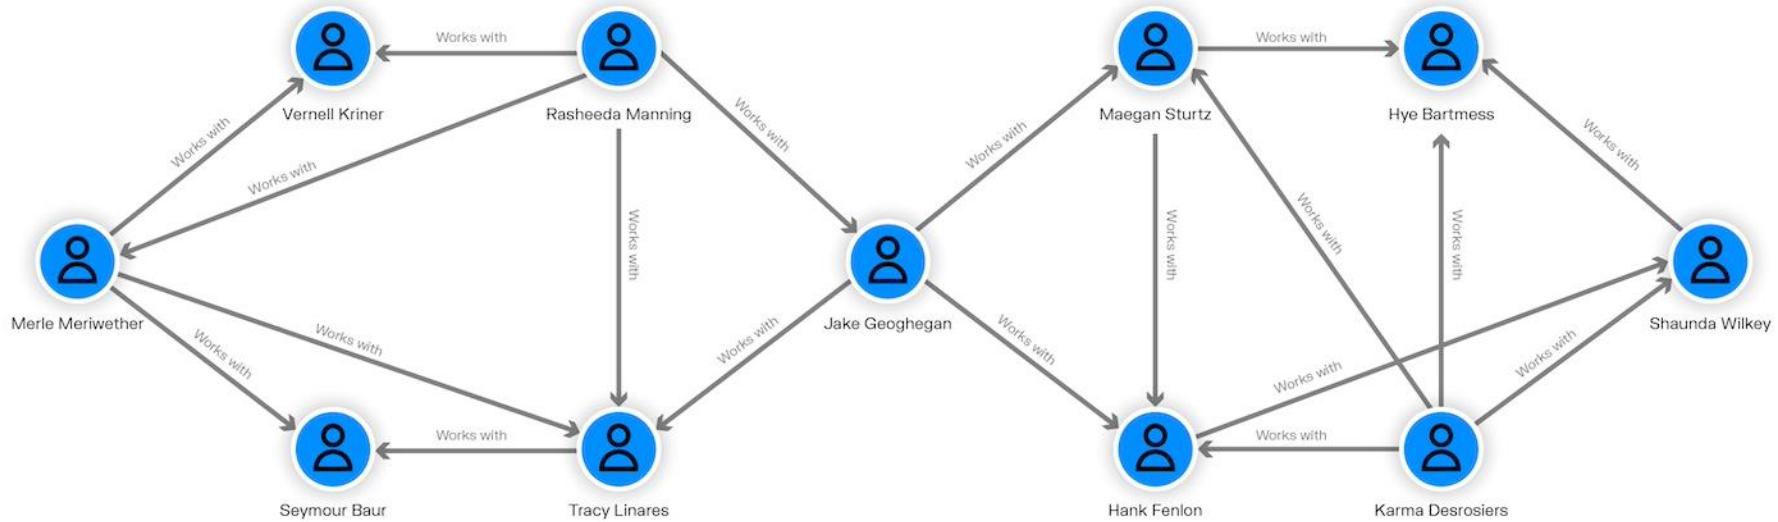
- Tasks as Series of Steps (Flow Chart Design)
- Access to Tools (Code, Web, Vector DB)
- Specialized LLMs per Step (Mix & Match Sizes/Strengths)

Common Patterns

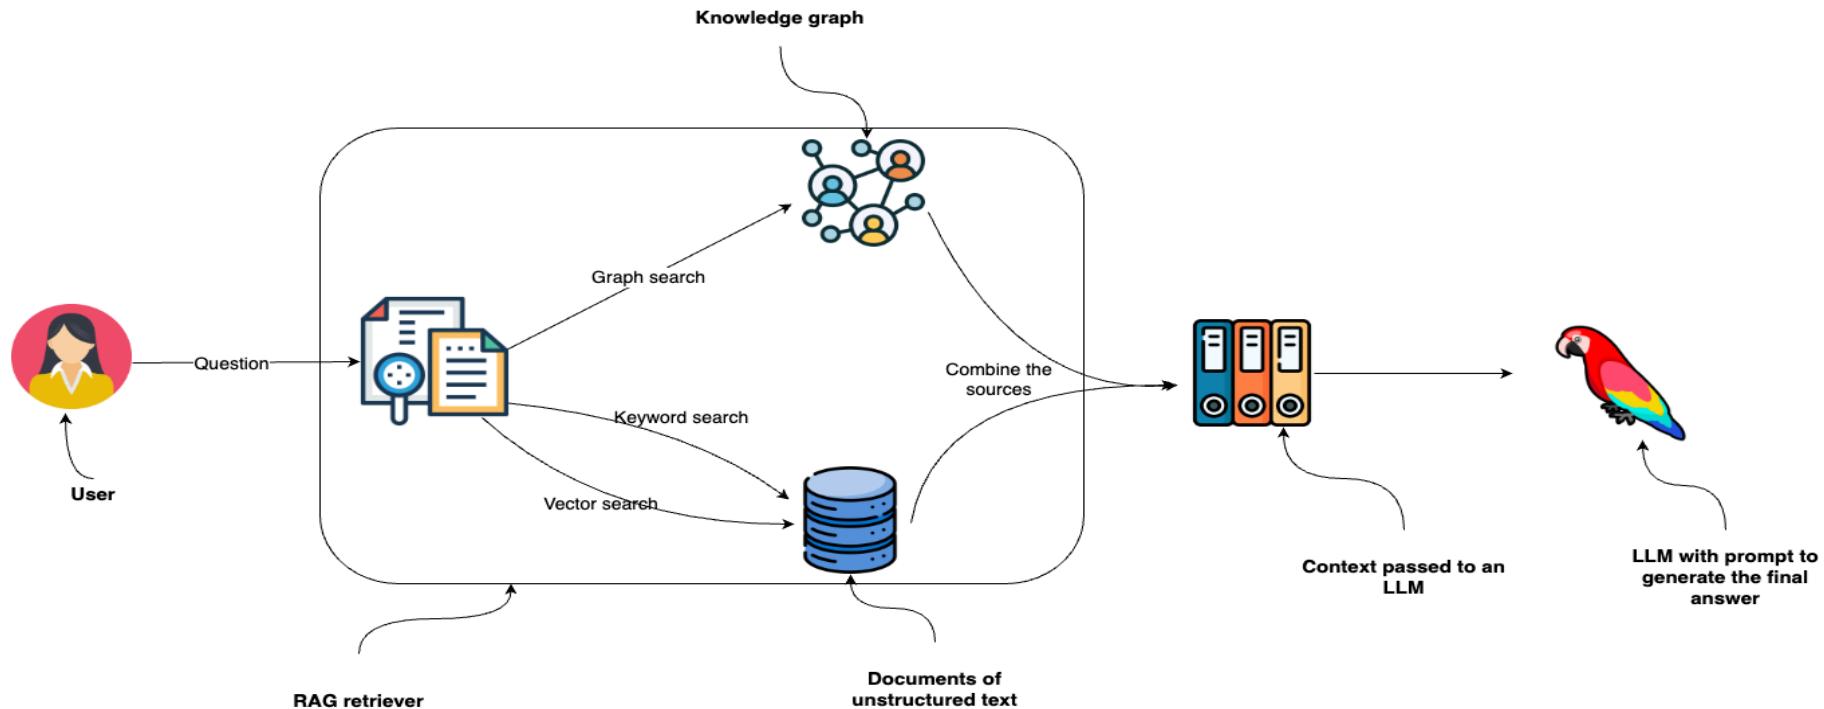


Knowledge Graphs

What is a Knowledge Graph



Knowledge Graphs + RAG

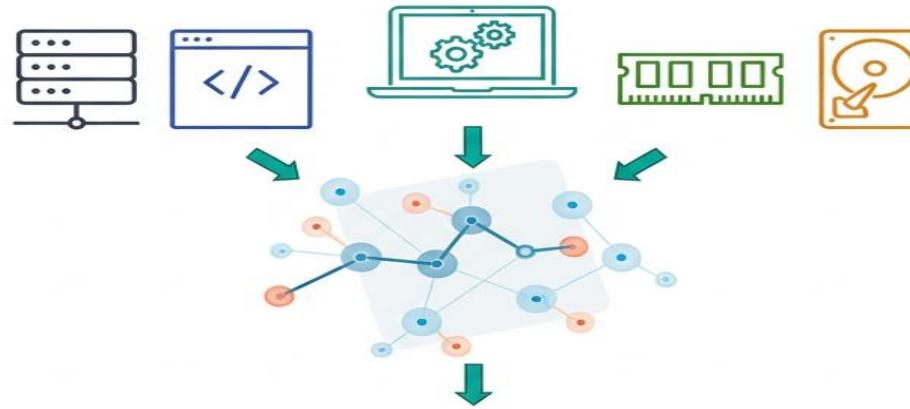


Graph-enabled RAG

Why is my new laptop running so slow?



**How
different
technical
details are
connected.**



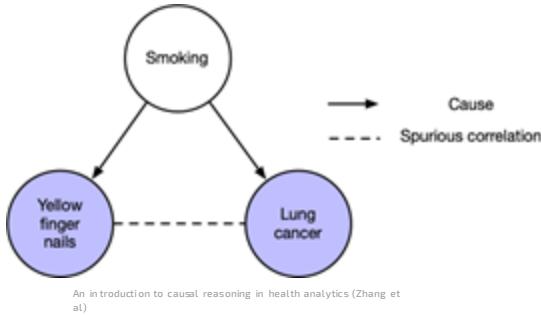
Your slow laptop performance, combined with high CPU usage from background apps and low available RAM, suggests insufficient memory for your workload. Consider upgrading RAM or closing unnecessary programs.



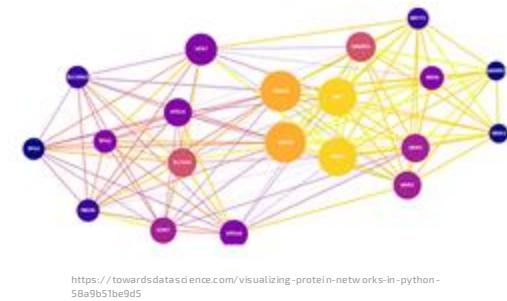
Why Graphs?



Knowledge Graph



Causal Graph



Interaction Graph

In many ways, graphs are the main modality of data we receive from nature.

– Petar Veličković

Why Graphs?

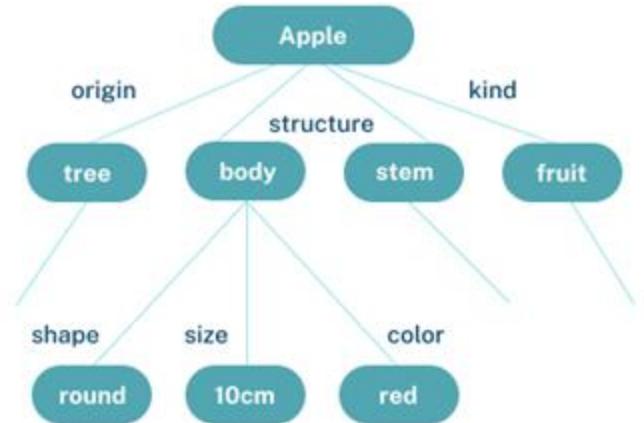


Human View of an Apple



Vector View of an Apple

Knowledge Graph View of an Apple



In many ways, graphs are the main modality of data we receive from nature.

- Petar Veličković

Why Graphs?

(Practically) all structured data can be converted into a graph & (almost) every graph can be considered a knowledge graph

Knowledge Graph

Causal Graph

Interaction Graph



In many ways, graphs are the main modality of data we receive from nature.

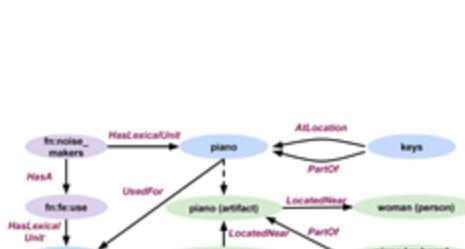
– Petar Veličković

What is a Knowledge Graph?

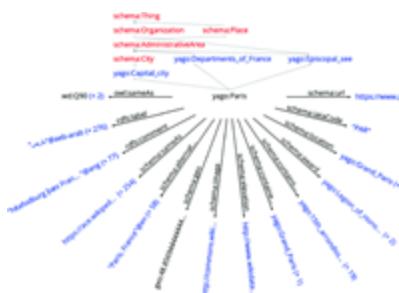
A KG consists of entities (**nodes**) that are connected by relationships (**edges**). Typically represented as list of (**subject, predicate, object**) triplets



WikiData
> 100 million items



CSKG
> 2 million items



Yago
> 10 million items



DrugBank
20 thousand items

What is a Knowledge Graph?

A KG consists of entities (**nodes**) that are connected by relationships (**edges**).
Typically represented as list of (**subject, predicate, object**) triplets

Create private KG



Financial Data



Medical Data

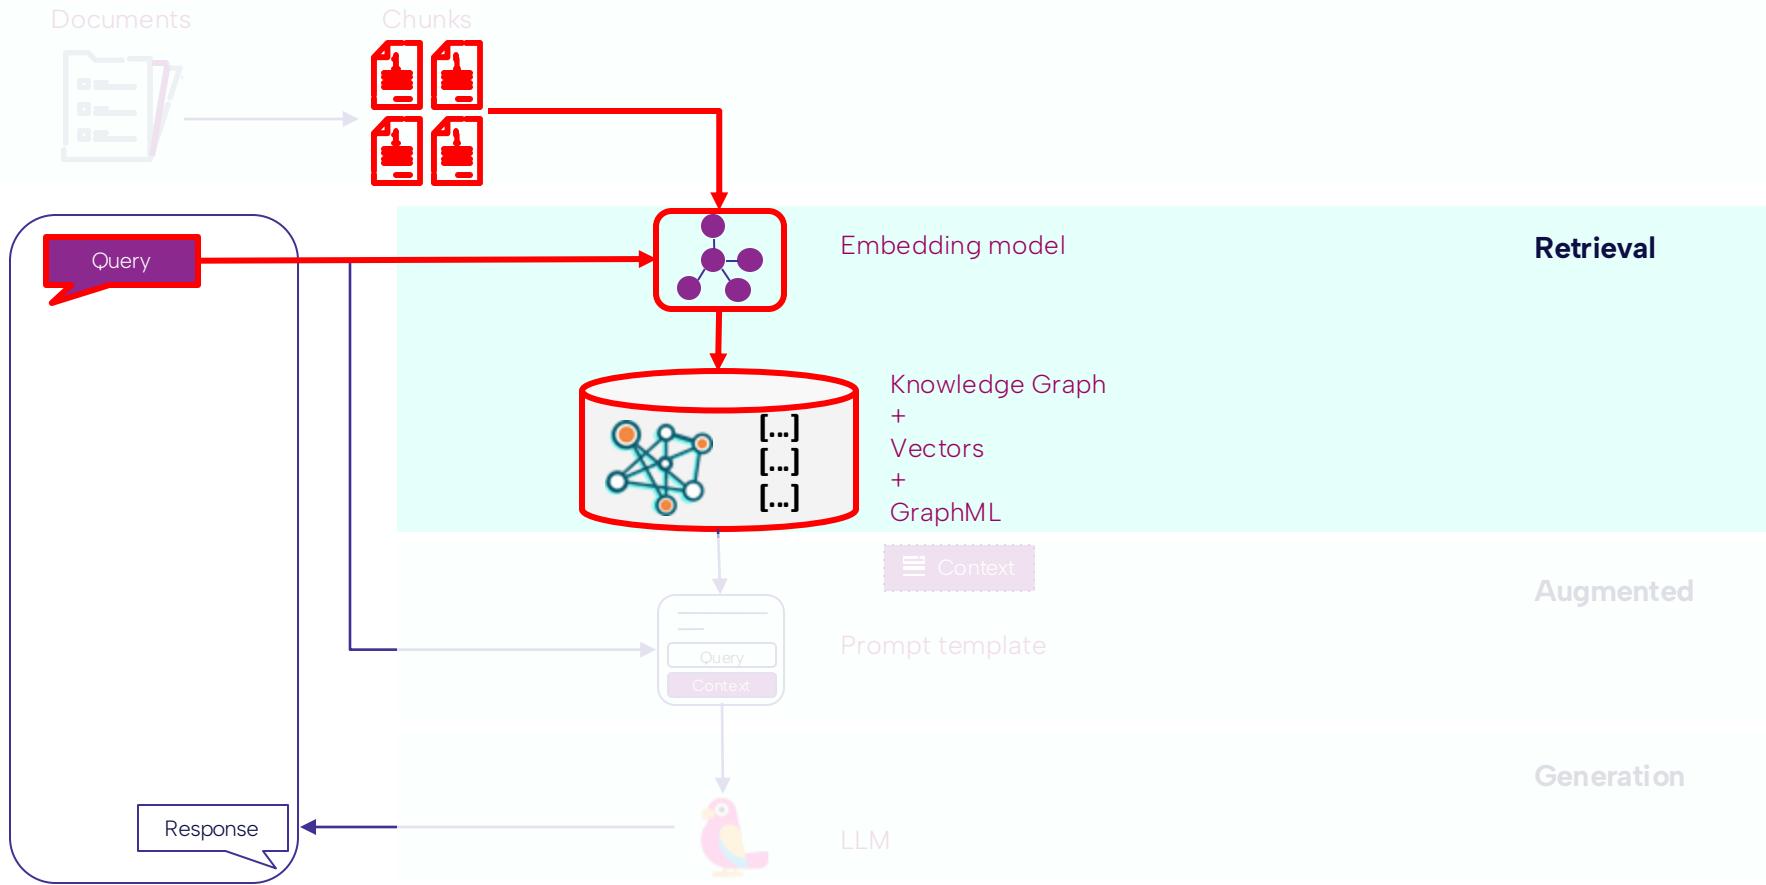


Vacation Planning



Reading List

GraphRAG



Thank you!

Engineering with GenAI
09.12.2025

